

## Homework 2

### Problem 1: Linear Regression (60 points)

#### 1.1 Print the shapes of these four objects. (5 points)

```
In [250]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 import mltools as ml
4
5 data = np.genfromtxt ( "curve80.txt", delimiter = None)
6 #The first column (data[:,0]) is the scalar feature value x;
7 X = data[:,0]
8
9 # code expects shape (M,N) so make sure it's 2-dimensional
10 X = np.atleast_2d(X).T
11
12 #The second column data[:,1] is the target value y for each example.
13 Y = data[:,1]
14 # split data set 75/25
15
16 Xtr,Xte,Ytr,Yte = ml.splitData(X,Y,0.75)
17
18 print ("Xtr = ", Xtr.shape)
19 print ("Xte = ", Xte.shape)
20 print ("Ytr = ", Ytr.shape)
21 print ("Yte = ", Yte.shape)
22
23
24
```

```
Xtr = (60, 1)
Xte = (20, 1)
Ytr = (60,)
Yte = (20,)
```

#### 1.2 Use the provided linearRegress class to create a linear regression predictor of y given x. You can plot the resulting function by simply evaluating the model at a large number of x values xs

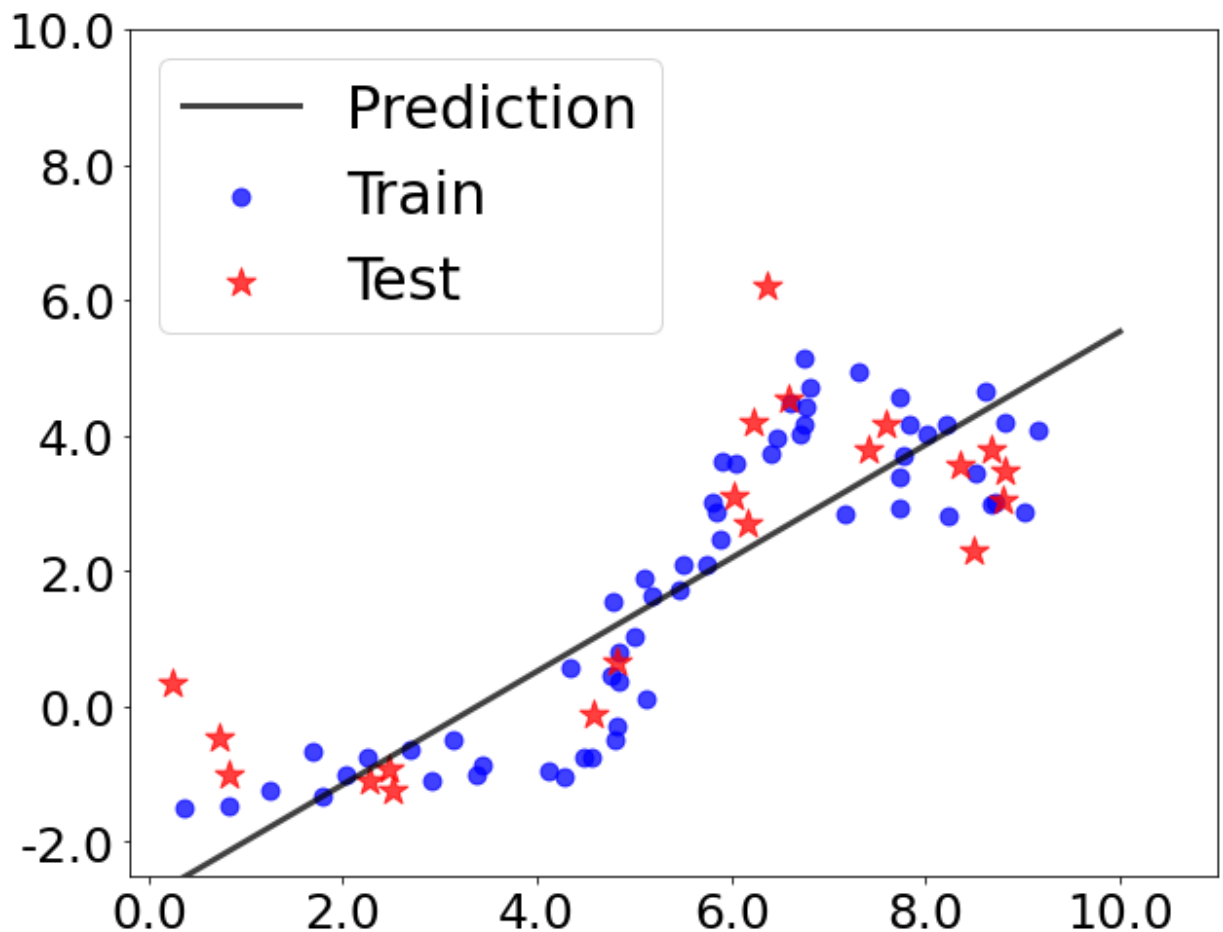
```
In [251]: 1 lr = ml.linear.linearRegress( Xtr, Ytr ) # create and train model
          2 xs = np.linspace(0,10,200) # densely sample possible x-values
          3 xs = xs[:,np.newaxis] # force "xs" to be an Mx1 matrix (expected by our
          4 ys = lr.predict( xs ) # make predictions at xs
```

**1.2.a : Plot the training data points along with your prediction function in a single plot. (10 points)**

```

In [252]: 1 # Plotting the data
2 f, ax = plt.subplots(1, 1, figsize=(10, 8))
3
4 ax.scatter(Xtr, Ytr, s=80, color='blue', alpha=0.75, label='Train')
5 ax.scatter(Xte, Yte, s=240, marker='*', color='red', alpha=0.75, label=
6
7 # Also plotting the regression line
8 ax.plot(xs, ys, lw=3, color='black', alpha=0.75, label='Prediction')
9
10 ax.set_xlim(-0.2, 11)
11 ax.set_ylim(-2.5, 10)
12 ax.set_xticklabels(ax.get_xticks(), fontsize=25)
13 ax.set_yticklabels(ax.get_yticks(), fontsize=25)
14
15 # Controlling the size of the legend and the location.
16 ax.legend(fontsize=30, loc=0)
17
18 plt.show()

```



```

In [5]: 1 print ( " Linear regression coefficients", lr.theta )

Linear regression coefficients [[-2.82765049  0.83606916]]

```

$$y = 0.836 + (-)2.8277x$$

## 1.2.c What is the mean squared error of the predictions on the training and test data? (10 points)

```
In [253]: 1 def MSE(y_true, y_hat):
2         y_true = y_true.reshape (-1, 1)
3         mse = np.sum ( (y_true - y_hat)**2 ) / len (y_true)
4
5         return mse
6 YtrHat = lr.theta[0][1] * Xtr + lr.theta[0][0]
7 YteHat = lr.theta[0][1] * Xte + lr.theta[0][0]
8
9 print('Mean Squared Error (training data) = ', MSE(Ytr, YtrHat))
10 print('Mean Squared Error (test data    ) = ', MSE(Yte, YteHat))
11
12
13
```

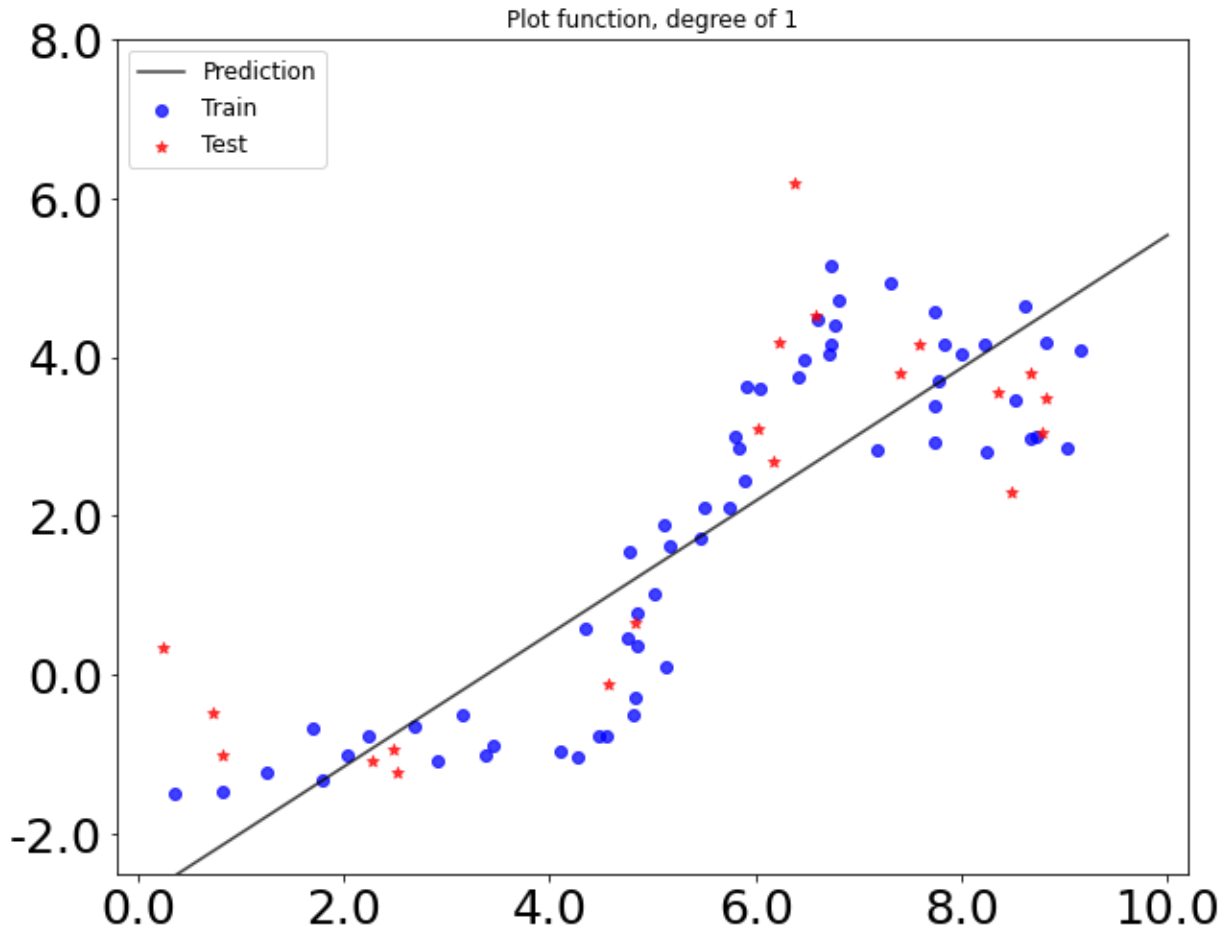
```
Mean Squared Error (training data) =  1.127711955609391
Mean Squared Error (test data    ) =  2.2423492030101246
```


## Promblem1.3


### 1.3.a : For each model, plot the learned prediction function $f(x)$ . (15 points)

In [254]:


```
1
2
3 degree = np.array([1,3,5,7,10,15,18])
4 for d in degree:
5
6     XtrP = ml.transforms.fpoly(Xtr, d, bias=False)
7     XtrP, params = ml.transforms.rescale(XtrP)
8
9     lr = ml.linear.linearRegress( XtrP, Ytr ) # create and train model
10
11
12     # Make sure you use the current space.
13     xs = np.linspace(0, 10, 200)
14     xs = np.atleast_2d(xs).T
15
16     xsP, _ = ml.transforms.rescale(ml.transforms.fpoly(xs, d, bias=False)
17
18     ys = lr.predict(xsP)
19
20     # draw graph
21     # Plotting the data
22     f, ax = plt.subplots(1, 1, figsize=(10, 8)) # size of graph
23     ax.scatter(Xtr, Ytr, color='blue', alpha=0.75, label='Train') #
24     ax.scatter(Xte, Yte, marker='*', color='red', alpha=0.75, label='Te
25     # Also plotting the regression line
26     ax.plot(xs, ys, color='black', alpha=0.75, label='Prediction')
27
28     plt.title( "Plot function, degree of {}".format(d))
29
30     ax.set_xlim(-0.2, 10.2)
31     ax.set_ylim(-2.5, 8)
32     ax.set_xticklabels(ax.get_xticks(), fontsize=25)
33     ax.set_yticklabels(ax.get_yticks(), fontsize=25)
34
35
36     ax.legend(fontsize=12, loc='upper left')
37
38     plt.show()
39
40
41
42
43
44
45
```


















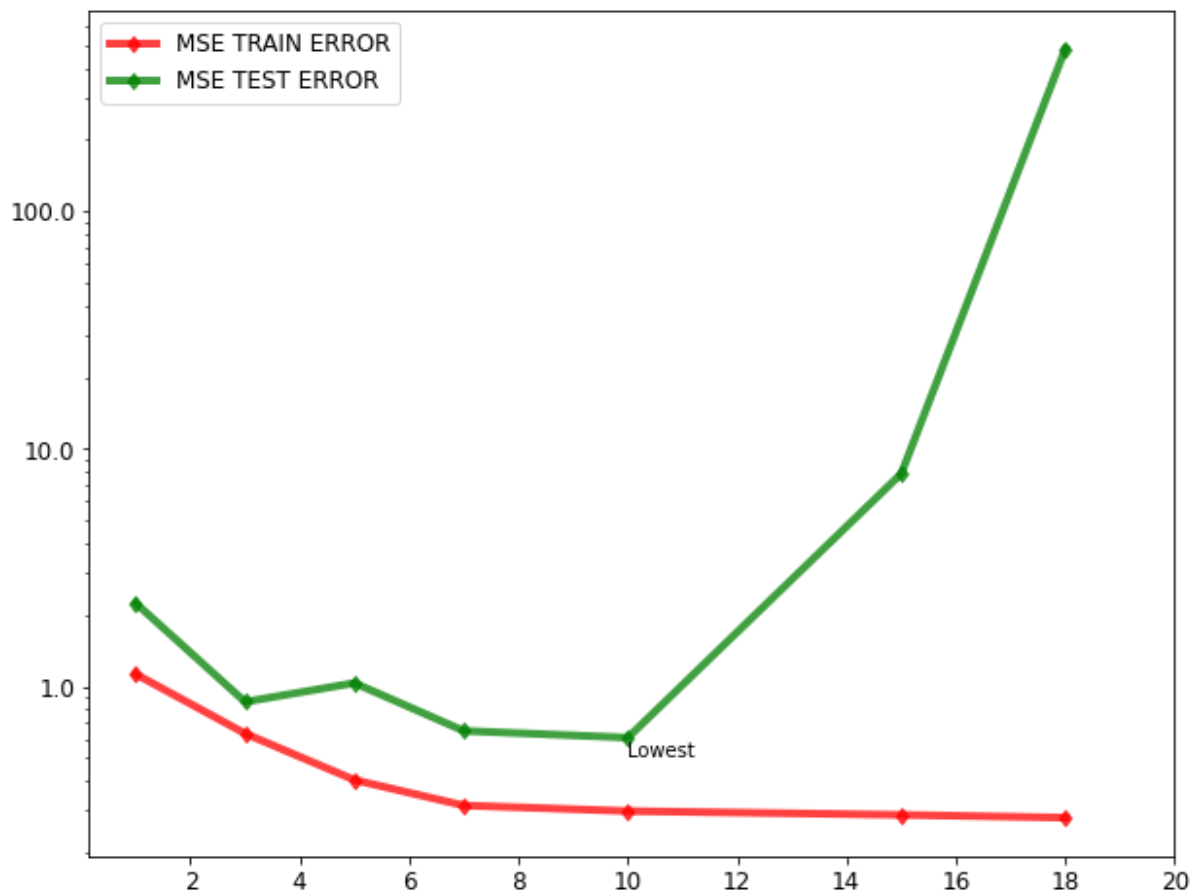
**1.2.b : Plot the training and test errors on a log scale ( semilogy ) as a function of the model degree. (10 points)**

```

In [256]: 1 degrees = np.array([1,3,5,7,10,15,18])
          2
          3 mse_train_error = np.zeros(degrees.shape[0])
          4 mse_test_error = np.zeros(degrees.shape[0])
          5 for i,degree in enumerate(degrees):
          6
          7     XtrP,params = ml.transforms.rescale(ml.transforms.fpoly(Xtr, degree
          8     lr = ml.linear.linearRegress(XtrP, Ytr)
          9     YtrHat = lr.predict(XtrP)
         10
         11     XteP,_ = ml.transforms.rescale(ml.transforms.fpoly(Xte, degree, bia
         12     YteHat = lr.predict(XteP)
         13
         14     mse_train_error[i] = MSE(Ytr, YtrHat)
         15     mse_test_error[i] = MSE(Yte, YteHat)
         16
         17
         18 fig, ax = plt.subplots(1, 1, figsize=(10, 8)) # Create axes for single
         19 # Plotting a line with markers where there's an actual x value.
         20 ax.semilogy(degrees, mse_train_error, lw=4, color = "red",marker='d', a
         21 ax.semilogy(degrees, mse_test_error, lw=4, color = "green", marker='d',
         22
         23 a = degrees
         24 b = mse_train_error
         25 c = mse_test_error
         26 t = Table([a, b, c], names=('degree', 'mse_train_error', 'mse_test_erro
         27 print ( t)
         28
         29 ax.text(10, 0.5090600748904027, 'Lowest' )
         30 ax.set_xticks(np.arange(2, 21, 2))
         31 ax.set_xticklabels(ax.get_xticks(), fontsize=12)
         32 ax.set_yticklabels(ax.get_yticks(), fontsize=12)
         33 ax.legend(fontsize=12, loc=0)
         34
         35 plt.show()

```

degree	mse_train_error	mse_test_error
1	1.1277119556093909	2.242349203010125
3	0.6339652063119635	0.8616114815449999
5	0.4042489464459056	1.0344190205632156
7	0.3156346739892996	0.6502246079670317
10	0.2989479796813433	0.6090600748904027
15	0.28817930796536423	7.863359085837317
18	0.28048505409585217	482.2803273735196



**1.2.c : What is the mean squared error of the predictions on the training and test data? (10 points)**

before degree of 10, our test data's performance improve. However, after the lowest point, degree of 15 and 18 are increasing which mean our test data is going worst.


**Problem 3: Statement of Collaboration (5 points)**

**1.3.1**

```

In [259]: 1 newDegree = [1, 3, 5, 7, 10, 15, 18]
          2 nFolds = 5; #given
          3
          4 J = np.zeros(nFolds)
          5 crsVal = []
          6 for i, degree in enumerate(newDegree):
          7     for iFold in range(nFolds):#given
          8
          9         Xti,Xvi,Yti,Yvi = ml.crossValidate(Xtr,Ytr,nFolds,iFold)
         10
         11         XtiP,params = ml.transforms.rescale(ml.transforms.fpoly(Xti, de
         12         learner = ml.linear.linearRegress(XtiP,Yti)
         13
         14
         15         XviP, _ = ml.transforms.rescale(ml.transforms.fpoly(Xvi, degree
         16
         17
         18         J[iFold] = learner.mse(XviP, Yvi)
         19
         20     crsVal.append(np.mean(J))
         21
         22 f, ax = plt.subplots(1, 1, figsize=(10, 8))    # size of graph
         23 ax.semilogy(newDegree,crsVal, lw=4, color = "blue",marker='d', alpha=0.
         24 #plt.semilogy(newDegree,mse_test_error, color = 'red')
         25 ax.legend(fontsize=15, loc=0)
         26 plt.show()

```



In [ ]:

1

**1.3.2 : How do the MSE estimates from five-fold cross-validation compare to the MSEs evaluated on the actual test data (Problem 1)? (5 points)**




```

In [187]: 1 a = degrees
          2 b = mse_test_error
          3 c = crsVal
          4 t = Table([a, b, c], names=('degree', 'mse_train_error', 'error'))
          5 print ( t)
          6
          7 f, ax = plt.subplots(1, 1, figsize=(10, 8))      # size of graph
          8 ax.semilogy(newDegree,crsVal, lw=4, color = "blue",marker='d', alpha=0.
          9 ax.semilogy(degrees, mse_test_error, lw=4, color = "green", marker='d',
         10 #plt.semilogy(newDegree,mse_test_error, color = 'red')
         11 ax.legend(fontsize=15, loc=0)
         12 plt.show()

```

degree	mse_train_error	error
1	2.242349203010125	1.2118626629641984
3	0.8616114815449999	0.7429005752051661
5	1.0344190205632156	0.5910703726406558
7	0.6502246079670317	0.7335637831345124
10	0.6090600748904027	0.7677056859101964
15	7.863359085837317	26.989609532127144
18	482.2803273735196	216818.07410494355



**1.3.3 : Which polynomial degree do you recommend based on five-fold cross-validation**

## error? (5 points)

**Based on five-fold cross-validation error, I recommend degree of 5. At this degree, we have the smallest error rate(0.5918).**

**1.3.4. For the degree that you picked in step 3, plot (with semilogy ) the cross-validation error as the number of folds is varied from nFolds = 2, 3, 4, 5, 6, 10, 12, 15. What pattern do you observe, and how do you explain why it occurs? (15 points)**

```
In [ ]: 1 nFolds = [2, 3, 4, 5, 6, 10, 12, 15]
2 myDegree = 5; #given
3
4 J = np.zeros(myDegree)
5 crsVal = []
6 for i, degree in enumerate(myDegree):
7     for iFold in range(nFolds):#given
8
9         Xti,Xvi,Yti,Yvi = ml.crossValidate(Xtr,Ytr,nFolds,iFold)
10
11         XtiP,params = ml.transforms.rescale(ml.transforms.fpoly(Xti, de
12         learner = ml.linear.linearRegress(XtiP,Yti)
13
14
15         XviP, _ = ml.transforms.rescale(ml.transforms.fpoly(Xvi, degree
16
17
18         J[iFold] = learner.mse(XviP, Yvi)
19
20         crsVal.append(np.mean(J))
21
22 f, ax = plt.subplots(1, 1, figsize=(10, 8))    # size of graph
23 ax.semilogy(newDegree,crsVal, lw=4, color = "blue",marker='d', alpha=0.
24 #plt.semilogy(newDegree,mse_test_error, color = 'red')
25 ax.legend(fontsize=15, loc=0)
26 plt.show()
```

## Problem 3: Statement of Collaboration (5 points)

- Piazza Question : question@222
- peter Park : Discuss about definition about MSE

